

Statistical Computing with R

Masters in Data Science 503 (S7&8)

Second Batch, SMS, TU, 2023

Shital Bhandary

Associate Professor

Statistics/Bio-statistics, Demography and Public Health Informatics

Patan Academy of Health Sciences, Lalitpur, Nepal

Faculty, Data Analysis and Decision Modeling, MBA, Pokhara University, Nepal

Faculty, FAIMER Fellowship in Health Professions Education, India/USA.

Review Preview (Unit 2, Part 2 & 3)

- Data wrangling
- Data munching
- Tidy data
- dplyr package and its use for data manipulation
- Database and R
- Reading database in R
- dbplyr package and its use for relational database manipulation with dplyr
- Text Mining

Data wrangling (Course book Chapter 9-16)

- Data wrangling is the art of getting your data into R in a useful form for visualization and modelling.
- Data wrangling is very important: without it you can't work with your own data! There are three main parts to data wrangling:
 - Import
 - Tidy
 - Transform

Import data in R:

- We have already covered this in the previous classes
- More here: <https://r4ds.had.co.nz/data-import.html>
- Reach this chapter well as there are some important import functions that are part of this course and may not have discussed so far
- We will discuss about reading “database” in the second part of this class

Tidy data in R

- Tidy data is a consistent way to organize your data in R.
- Getting your/our data into this format requires some upfront work, but that work pays off in the long term.
- Once you/we have tidy data and the tidy tools provided by packages in the **tidyverse**, you will spend much less time **munging/cleaning data** from one representation to another, allowing you to spend more time on the analytic questions at hand.
- Tidy data in tidyverse packages are stored as “tibble”

Let us see what is “tibble” first:

<https://r4ds.had.co.nz/tibbles.html>

- The variant of the data frame used by “tidiverse” is called: **tibble**.
- Tibbles are data frames, but they tweak some older behaviours to make life a little easier.
- R is an old language, and some things that were useful 10 or 20 years ago now get in your way.
- It's difficult to change base R without breaking existing code, so **most innovation occurs in packages**.
- The **tibble** package provides opinionated data frames that make working in the **tidyverse** a little easier.
- It's particularly **useful for large datasets** because it only prints the first few rows.

Note:

- All the functions of “tidyverse” package works fast with tibble so it will be wise to say that data frame/s should be converted to tibble before using functions of “tidyverse” package
- However, most of the packages of the “tidyverse” super package works well with the data frame too!
- There are two main differences in the usage of a data frame vs a tibble: printing, and subsetting. <https://posit.co/blog/tibble-1-0-0/>

There are three interrelated rules which make a dataset tidy:

- Each variable must have its own column.
- Each observation must have its own row.
- Each value must have its own cell.

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	216766	128042583

variables

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	216766	128042583

observations

country	year	cases	population
Afghanistan	99	745	19987071
Afghanistan	00	2666	20595360
Brazil	99	37737	172006362
Brazil	00	80488	174504898
China	99	212258	1272915272
China	00	216766	128042583

values

Example: Which one is “tidy”? Why?

#Creating tibble:

```
table1 <- tibble(  
  country = c("Afghanistan", "Afghanistan", "Brazil",  
    "Brazil", "China", "China"),  
  year = c(1999, 2000, 1999, 2000, 1999, 2000),  
  cases = c(745, 2666, 37737, 80488, 212258, 213766),  
  population = c(19987071, 20595360, 172006362,  
    174504898, 1272915272, 1280428583)  
)
```

Data frame to tibble: `as_tibble(data_frame)`

Tibble to data frame: `as.data.frame(tibble_data)`

- table1
- A tibble: 6 x 4
- | | year | country | cases | population |
|-----|-------|-------------|--------|------------|
| | <dbl> | <chr> | <dbl> | <dbl> |
| • 1 | 1999 | Afghanistan | 745 | 19987071 |
| • 2 | 2000 | Afghanistan | 2666 | 20595360 |
| • 3 | 1999 | Brazil | 37737 | 172006362 |
| • 4 | 2000 | Brazil | 80488 | 174504898 |
| • 5 | 1999 | China | 212258 | 1272915272 |
| • 6 | 2000 | China | 213766 | 1280428583 |

- **dbl = duple instead of number in “tibble”!**

Example: Which one is “tidy”? Why?

- table2
- #> # A tibble: 12 × 4
- #> country year type count
- #> <chr> <int> <chr> <int>
- #> 1 Afghanistan 1999 cases 745
- #> 2 Afghanistan 1999 population 19987071
- #> 3 Afghanistan 2000 cases 2666
- #> 4 Afghanistan 2000 population 20595360
- #> 5 Brazil 1999 cases 37737
- #> 6 Brazil 1999 population 172006362
- #> # ... with 6 more rows

- table3
- #> # A tibble: 6 × 3
- #> country year rate
- #> * <chr> <int> <chr>
- #> 1 Afghanistan 1999 745/19987071
- #> 2 Afghanistan 2000 2666/20595360
- #> 3 Brazil 1999 37737/172006362
- #> 4 Brazil 2000 80488/174504898
- #> 5 China 1999 212258/1272915272
- #> 6 China 2000 213766/1280428583

Example: Which one is “tidy”? Why?

Spread across two tibbles

table4a # cases

- #> # A tibble: 3 × 3
- #> country `1999` `2000`
- #> * <chr> <int> <int>
- #> 1 Afghanistan 745 2666
- #> 2 Brazil 37737 80488
- #> 3 China 212258 213766

table4b # population

- #> # A tibble: 3 × 3
- #> country `1999` `2000`
- #> * <chr> <int> <int>
- #> 1 Afghanistan 19987071 20595360
- #> 2 Brazil 172006362 174504898
- #> 3 China 1272915272 1280428583

Why ensure that your data is tidy? There are two main advantages:

- There's a general advantage to picking one consistent way of storing data. If you have a consistent data structure, it's easier to learn the tools that work with it because they have an underlying uniformity.
- There's a specific advantage to placing variables in columns because it allows R's vectorized nature to shine.
- **dplyr**, **ggplot2**, and all the other packages in the **tidyverse** are designed to work with tidy data.

Tidy data: Pivoting – Longer to wider

- table2
- #> # A tibble: 12 × 4
- #> country year type count
- #> <chr> <int> <chr> <int>
- #> 1 Afghanistan 1999 cases 745
- #> 2 Afghanistan 1999 population 19987071
- #> 3 Afghanistan 2000 cases 2666
- #> 4 Afghanistan 2000 population 20595360
- #> 5 Brazil 1999 cases 37737
- #> 6 Brazil 1999 population 172006362
- #> # ... with 6 more rows

```
table2 %>%  
  pivot_wider(names_from = type, values_from =  
    count)
```

country	year	key	value
Afghanistan	1999	cases	745
Afghanistan	1999	population	19987071
Afghanistan	2000	cases	2666
Afghanistan	2000	population	20595360
Brazil	1999	cases	37737
Brazil	1999	population	172006362
Brazil	2000	cases	80488
Brazil	2000	population	174504898
China	1999	cases	212258
China	1999	population	1272915272
China	2000	cases	213766
China	2000	population	1280428583

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	213766	1280428583

table2

Figure 12.3: Pivoting table2 into a wider, tidy form.

Tidy data: Pivoting – Wider to Longer

- table4a
- #> # A tibble: 3 × 3
- #> country `1999` `2000`
- #> * <chr> <int> <int>
- #> 1 Afghanistan 745 2666
- #> 2 Brazil 37737 80488
- #> 3 China 212258 213766

```
table4a %>%  
  pivot_longer(c(`1999`, `2000`), names_to =  
    "year", values_to = "cases")
```

country	year	cases
Afghanistan	1999	745
Afghanistan	2000	2666
Brazil	1999	37737
Brazil	2000	80488
China	1999	212258
China	2000	213766

country	1999	2000
Afghanistan	745	2666
Brazil	37737	80488
China	212258	213766

table4

Figure 12.2: Pivoting table4 into a longer, tidy form.

Tidy data: Separate

- table3
- #> # A tibble: 6 × 3
- #> country year rate
- #> * <chr> <int> <chr>
- #> 1 Afghanistan 1999 745/19987071
- #> 2 Afghanistan 2000 2666/20595360
- #> 3 Brazil 1999 37737/172006362
- #> 4 Brazil 2000 80488/174504898
- #> 5 China 1999 212258/1272915272
- #> 6 China 2000 213766/1280428583

```
table3 %>%  
  separate(rate, into = c("cases", "population"))  
OR  
table3 %>%  
  separate(rate, into = c("cases", "population"), sep = "/")
```

The diagram illustrates the transformation of a data table from a wide format to a tidy format. On the left, a table labeled 'table3' has three columns: 'country', 'year', and 'rate'. The 'rate' column contains strings in the format 'cases/population'. On the right, the resulting tidy table has four columns: 'country', 'year', 'cases', and 'population'. A curved arrow points from the 'rate' column of the original table to the 'cases' and 'population' columns of the new table, indicating the separation of the data.

country	year	rate
Afghanistan	1999	745 / 19987071
Afghanistan	2000	2666 / 20595360
Brazil	1999	37737 / 172006362
Brazil	2000	80488 / 174504898
China	1999	212258 / 1272915272
China	2000	213766 / 1280428583

table3

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	213766	1280428583

Figure 12.4: Separating table3 makes it tidy

Tidy data: Unite

- **unite()** is the inverse of **separate()**: it combines multiple columns into a single column.
- You'll need it much less frequently than **separate()**, but it's still a useful tool to have in your back pocket.

```
table5 %>%  
  unite(new, century, year)  
OR  
table5 %>%  
  unite(new, century, year, sep = "")
```

country	year	rate
Afghanistan	1999	745 / 19987071
Afghanistan	2000	2666 / 20595360
Brazil	1999	37737 / 172006362
Brazil	2000	80488 / 174504898
China	1999	212258 / 1272915272
China	2000	213766 / 1280428583

country	century	year	rate
Afghanistan	19	99	745 / 19987071
Afghanistan	20	0	2666 / 20595360
Brazil	19	99	37737 / 172006362
Brazil	20	0	80488 / 174504898
China	19	99	212258 / 1272915272
China	20	0	213766 / 1280428583

table6

Figure 12.5: Uniting table5 makes it tidy

Missing values

- Changing the representation of a dataset brings up an important subtlety of missing values.
- Surprisingly, a value can be missing in one of two possible ways:
 - **Explicitly**, i.e. flagged with NA.
 - **Implicitly**, i.e. simply not present in the data.

Missing values: Example

#Create a tibble with missing values:

```
stocks <- tibble(  
  year = c(2015, 2015, 2015, 2015, 2016, 2016, 2016),  
  qtr  = c( 1,  2,  3,  4,  2,  3,  4),  
  return = c(1.88, 0.59, 0.35, NA, 0.92, 0.17, 2.66)  
)
```

Missing values: Example

There are two missing values in this dataset:

- The return for the fourth quarter of 2015 is **explicitly missing**, because the cell where its value should be instead contains NA.
- The return for the first quarter of 2016 is **implicitly missing**, because it simply does not appear in the dataset.

Missing values: Example

- `stocks %>%`
- `pivot_wider(names_from = year, values_from = return)`
- `#> # A tibble: 4 × 3`
- `#>`

	qtr	`2015`	`2016`
<code>#></code>	<code><dbl></code>	<code><dbl></code>	<code><dbl></code>
<code>#> 1</code>	1	1.88	NA
<code>#> 2</code>	2	0.59	0.92
<code>#> 3</code>	3	0.35	0.17
<code>#> 4</code>	4	NA	2.66
- `#>`
- `#>`
- `#> 1`
- `#> 2`
- `#> 3`
- `#> 4`

Missing values: What will happen now?

- `stocks %>%`

```
  pivot_wider(names_from = year, values_from = return) %>%
```

```
  pivot_longer(
```

```
    cols = c(`2015`, `2016`),
```

```
    names_to = "year",
```

```
    values_to = "return",
```

```
    values_drop_na = TRUE
```

```
)
```

Missing values:

We can use “complete” command!

- `stocks %>%`
- `complete(year, qtr)`
- `#> # A tibble: 8 × 3`
- `#> year qtr return`
- `#> <dbl> <dbl> <dbl>`
- `#> 1 2015 1 1.88`
- `#> 2 2015 2 0.59`
- `#> 3 2015 3 0.35`
- `#> 4 2015 4 NA`
- `#> 5 2016 1 NA`
- `#> 6 2016 2 0.92`
- `#> # ... with 2 more rows`

Missing values: Another example (tibble by row or tribble!)

- `treatment <- tribble(
 ~ person, ~ treatment, ~response,
 "Derrick Whitmore", 1, 7,
 NA, 2, 10,
 NA, 3, 9,
 "Katherine Burke", 1, 4
)`
 - `treatment`

Missing values: fill() for another example

- treatment %>%
- fill(person) **# “tidyr” package is required here!**
- #> # A tibble: 4 × 3
- #> person treatment response
- #> <chr> <dbl> <dbl>
- #> 1 Derrick Whitmore 1 7
- #> 2 Derrick Whitmore 2 10
- #> 3 Derrick Whitmore 3 9
- #> 4 Katherine Burke 1 4

Question/Queries?

Transform/manipulate data with “dplyr”

- To learn five key dplyr functions that allow you to solve the vast majority of your data manipulation challenges:
 - Pick observations by their values (**filter()**).
 - Reorder the rows (**arrange()**).
 - Pick variables by their names (**select()**).
 - Create new variables with functions of existing variables (**mutate()**).
 - Collapse many values down to a single summary (**summarise()**).
- These can all be used in conjunction with **group_by()** which changes the scope of each function from operating on the entire dataset to operating on it group-by-group.

Data manipulation with “dplyr”

- These six functions provide the verbs for a language of data manipulation.
- All verbs work similarly:
 - The first argument is a data frame.
 - The subsequent arguments describe what to do with the data frame, using the variable names (without quotes).
 - The result is a new data frame.
- Together these properties make it easy to chain together multiple simple steps to achieve a complex result.

Let's use them with nycflighst13 data

- library(dplyr)
- library(nycflights13)
- flights
- #> # A tibble: 336,776 × 19
- #> year month day dep_time sched_dep...¹ dep_d...² arr_t...³ sched...⁴ arr_d...⁵ carrier
- #> <int> <int> <int> <int> <int> <dbl> <int> <int> <dbl> <chr>
- #> 1 2013 1 1 517 515 2 830 819 11 UA
- #> 2 2013 1 1 533 529 4 850 830 20 UA
- #> 3 2013 1 1 542 540 2 923 850 33 AA
- #> 4 2013 1 1 544 545 -1 1004 1022 -18 B6
- #> 5 2013 1 1 554 600 -6 812 837 -25 DL
- #> 6 2013 1 1 554 558 -4 740 728 12 UA
- #> # ... with 336,770 more rows, 9 more variables

Filter: What will happen?

- **filter(flights, month == 1, day == 1)**
- #> # **A tibble: 842 × 19**
- #> year month day dep_time sched_dep...¹ dep_d...² arr_t...³ sched...⁴
 arr_d...⁵ carrier
- #> <int> <int> <int> <int> <int> <dbl> <int> <int> <dbl> <chr>
- #> 1 2013 1 1 517 515 2 830 819 11 UA
- #> 2 2013 1 1 533 529 4 850 830 20 UA
- #> 3 2013 1 1 542 540 2 923 850 33 AA
- #> 4 2013 1 1 544 545 -1 1004 1022 -18 B6
- #> 5 2013 1 1 554 600 -6 812 837 -25 DL
- #> 6 2013 1 1 554 558 -4 740 728 12 UA
- #> # ... with 836 more rows, 9 more variables

Are these better?

- `jan1 <- filter(flights, month == 1, day == 1)`
- `(jan1 <- filter(flights, month == 1, day == 1))`

- `dec25 <- filter(flights, month == 12, day == 25)`
- `(dec25 <- filter(flights, month == 12, day == 25))`

- `filter(flights, month = 1)` `#Why error?`
- `filter(flights, month == 1)` `#Works now? Why?`

More with filter:

- `filter(flights, month == 11 | month == 12)` #What?
- `filter(flights, month == 11 | 12)` #Works?
- `nov_dec <- filter(flights, month %in% c(11, 12))` #Works?
- De Morgan's Law:
 - `filter(flights, !(arr_delay > 120 | dep_delay > 120))` #Works?
 - `filter(flights, arr_delay <= 120, dep_delay <= 120)` #Works?

Arrange: Example

- `arrange(flights, year, month, day)`
- `#> # A tibble: 336,776 × 19`
- `#> year month day dep_time sched_dep...1 dep_d...2 arr_t...3 sched...4 arr_d...5
carrier`
- `#> <int> <int> <int> <int> <int> <dbl> <int> <int> <dbl> <chr>`
- `#> 1 2013 1 1 517 515 2 830 819 11 UA`
- `#> 2 2013 1 1 533 529 4 850 830 20 UA`
- `#> 3 2013 1 1 542 540 2 923 850 33 AA`
- `#> 4 2013 1 1 544 545 -1 1004 1022 -18 B6`
- `#> 5 2013 1 1 554 600 -6 812 837 -25 DL`
- `#> 6 2013 1 1 554 558 -4 740 728 12 UA`
- `#> # ... with 336,770 more rows, 9 more variables`

What will happen now?

- Arrange will sort the data in ascending order
- `arrange(flights, desc(dep_delay))`
- Use `desc()` to re-order by a column in descending order
- Missing values are always sorted at the end

Select: Example

- # Select columns by name
- `select(flights, year, month, day)`
- `#> # A tibble: 336,776 × 3`
- `#> year month day`
- `#> <int> <int> <int>`
- `#> 1 2013 1 1`
- `#> 2 2013 1 1`
- `#> 3 2013 1 1`
- `#> 4 2013 1 1`
- `#> 5 2013 1 1`
- `#> 6 2013 1 1`
- `#> # ... with 336,770 more rows`

- # Select all columns between year and day (inclusive)
- `select(flights, year:day)`
- `#> # A tibble: 336,776 × 3`
- `#> year month day`
- `#> <int> <int> <int>`
- `#> 1 2013 1 1`
- `#> 2 2013 1 1`
- `#> 3 2013 1 1`
- `#> 4 2013 1 1`
- `#> 5 2013 1 1`
- `#> 6 2013 1 1`
- `#> # ... with 336,770 more rows`

Select: “except” example

- # Select all columns except those from year to day (inclusive)
- `select(flights, -(year:day))`
- #> # A tibble: 336,776 × 16
- #> dep_time sched...¹ dep_d...² arr_t...³ sched...⁴ arr_d...⁵ carrier flight tailnum origin
- #> <int> <int> <dbl> <int> <int> <dbl> <chr> <int> <chr> <chr>
- #> 1 517 515 2 830 819 11 UA 1545 N14228 EWR
- #> 2 533 529 4 850 830 20 UA 1714 N24211 LGA
- #> 3 542 540 2 923 850 33 AA 1141 N619AA JFK
- #> 4 544 545 -1 1004 1022 -18 B6 725 N804JB JFK
- #> 5 554 600 -6 812 837 -25 DL 461 N668DN LGA
- #> 6 554 558 -4 740 728 12 UA 1696 N39463 EWR
- #> # ... with 336,770 more rows, 6 more variables

Select: More

- There are a number of helper functions you can use within `select()`:
- `starts_with("abc")`: matches names that begin with “abc”.
- `ends_with("xyz")`: matches names that end with “xyz”.
- `contains("ijk")`: matches names that contain “ijk”.
- `matches("(.)\\1")`: selects variables that match a **regular expression**.
- This one matches any variables that contain repeated characters.
- `num_range("x", 1:3)`: matches x1, x2 and x3.
- See `?select` for more details.

More on regular expression are available here: <https://cran.r-project.org/web/packages/stringr/vignettes/regular-expressions.html>

Note:

- `select()` can be used to rename variables, but it's rarely useful because it drops all of the variables not explicitly mentioned.
- Instead, use `rename()`, which is a variant of `select()` that keeps all the variables that aren't explicitly mentioned
- `rename(flights, tail_num = tailnum)`
- Another option is to use `select()` in conjunction with the `everything()` helper.
- This is useful if you have a handful of variables you'd like to move to the start of the data frame.
- `select(flights, time_hour, air_time, everything())`

Mutate: Example

- Besides selecting sets of existing columns, it's often useful to add new columns that are functions of existing columns.
- That's the job of `mutate()`.
- `mutate()` always adds new columns at the end of your dataset so we'll start by creating a narrower dataset so we can see the new variables.

```
#Adding variables in flights_sml:
flights_sml <- select(flights,
  year:day,
  ends_with("delay"),
  distance,
  air_time
)
mutate(flights_sml,
  gain = dep_delay - arr_delay,
  speed = distance / air_time * 60
)
```

Mutate: Example

- Besides selecting sets of existing columns, it's often useful to add new columns that are functions of existing columns.
- That's the job of `mutate()`.
- `mutate()` always adds new columns at the end of your dataset so we'll start by creating a narrower dataset so we can see the new variables.

#Adding one more variable:

```
mutate(flights_sml,  
       gain = dep_delay - arr_delay,  
       hours = air_time / 60,  
       gain_per_hour = gain / hours  
)
```

#Note that you/we can refer to columns that you've just created

Transmute and other useful creation functions

More@ <https://r4ds.had.co.nz/transform.html>

- If you only want to keep the new variables, use `transmute()`
 - Arithmetic operators: `+`, `-`, `*`, `/`, `^`
 - Modular arithmetic: `%/%` (integer division) and `%%` (remainder)
 - Use: Compute hour and minute from `dep_time` with:
 - `transmute(flights, dep_time, hour = dep_time %/% 100, minute = dep_time %% 100)`
- ```
transmute(flights,
 gain = dep_delay - arr_delay,
 hours = air_time / 60,
 gain_per_hour = gain / hours
)
```



# Summarise: Works best for group summaries

- `summarise(flights, delay = mean(dep_delay, na.rm = TRUE))`
- `#> # A tibble: 1 × 1`
- `#> delay`
- `#> <dbl>`
- `#> 1 12.6`
- `by_day <- group_by(flights, year, month, day)`
- `summarise(by_day, delay = mean(dep_delay, na.rm = TRUE))`
- `#> # A tibble: 365 × 4`
- `#> # Groups: year, month [12]`
- `#> year month day delay`
- `#> <int> <int> <int> <dbl>`
- `#> 1 2013 1 1 11.5`
- `#> 2 2013 1 2 13.9`
- `#> 3 2013 1 3 11.0`
- `#> 4 2013 1 4 8.95`
- `#> 5 2013 1 5 5.73`
- `#> 6 2013 1 6 7.15`
- `#> # ... with 359 more rows`

# Multiple operations: pipes

```
#What will happen?
delays <- flights %>%
 group_by(dest) %>%
 summarise(
 count = n(),
 dist = mean(distance, na.rm =
TRUE),
 delay = mean(arr_delay, na.rm =
TRUE)
) %>%
 filter(count > 20, dest != "HNL")
```

```
#What will happen?
flights %>%
 group_by(year, month, day) %>%
 summarise(mean =
mean(dep_delay))

#And now?
flights %>%
 group_by(year, month, day) %>%
 summarise(mean =
mean(dep_delay, na.rm = TRUE))
```

# How to remove cancelled flights?

## And, get summaries by groups!

```
not_cancelled <- flights %>%
 filter(!is.na(dep_delay),
 !is.na(arr_delay))

not_cancelled %>%
 group_by(year, month, day) %>%
 summarise(mean =
 mean(dep_delay))
```

- #> # A tibble: 365 × 4
- #> # Groups: year, month [12]
- #> year month day mean
- #> <int> <int> <int> <dbl>
- #> 1 2013 1 1 11.4
- #> 2 2013 1 2 13.7
- #> 3 2013 1 3 10.9
- #> 4 2013 1 4 8.97
- #> 5 2013 1 5 5.73
- #> 6 2013 1 6 7.15
- #> # ... with 359 more rows

# Counts: Example

- Whenever you do any aggregation, it's always a good idea to include either a count (`n()`), or a count of non-missing values (`sum(!is.na(x))`).
- That way you can check that you're not drawing conclusions based on very small amounts of data.

# What happens now?

```
delays <- not_cancelled %>%
 group_by(tailnum) %>%
 summarise(
 delay = mean(arr_delay)
)

hist(delays$delay)
```

# What happens now?

```
delays <- not_cancelled %>%
 group_by(tailnum) %>%
 summarise(
 delay = mean(arr_delay,
na.rm = TRUE),
 n = n()
)
```

# Plots

```
hist(delays$n)
```

```
hist(delays$delay)
```

```
plot(delays$n, delays$delay)
```

# Useful summary functions:

<https://r4ds.had.co.nz/transform.html>

# When do the first and last flights  
leave each day?

```
not_cancelled %>%
```

```
 group_by(year, month, day) %>%
```

```
 summarise(
```

```
 first = min(dep_time),
```

```
 last = max(dep_time)
```

```
)
```

- # Why is distance to some  
destinations more variable than  
to others?

```
not_cancelled %>%
```

```
 group_by(dest) %>%
```

```
 summarise(distance_sd =
```

```
 sd(distance)) %>%
```

```
 arrange(desc(distance_sd))
```

# Useful summary functions:

<https://r4ds.had.co.nz/transform.html>

# Which destinations have the most carriers?

```
not_cancelled %>%
 group_by(dest) %>%
 summarise(carriers =
n_distinct(carrier)) %>%
 arrange(desc(carriers))
```

- # How many flights left before 5am? (these usually indicate delayed flights from the previous day)

```
not_cancelled %>%
 group_by(year, month, day) %>%
 summarise(n_early =
sum(dep_time < 500))
```

# Useful summary functions:

<https://r4ds.had.co.nz/transform.html>

# What proportion of flights are  
delayed by more than an hour?

```
not_cancelled %>%
 group_by(year, month, day) %>%
 summarise(hour_prop =
 mean(arr_delay > 60))
```

#Find all groups bigger than a  
threshold:

```
popular_dests <- flights %>%
 group_by(dest) %>%
 filter(n() > 365)
popular_dests
```



# Popular destination: head and tail

(Are these results VALID?)

- `head(popular_dests$dest)`
- `[1] "IAH" "IAH" "MIA" "BQN"`  
`"ATL" "ORD"`

- IAH = Texas
- MIA = Miami
- BQN = Puerto Rico
- ATL = Atlanta
- ORD = Chicago

- `tail(popular_dests$dest)`
- `[1] "BNA" "DCA" "SYR" "BNA"`  
`"CLE" "RDU"`

- BNA = Nashville
- DCA = Washington (Reagan Nat.)
- SYR = New York (Syracuse)
- CLE = Cleveland
- RDU = North Carolina

# Bonus: dplyr “slice” function with examples

<https://dplyr.tidyverse.org/reference/slice.html>

#What will happen?

`flights %>% slice(1L)`

`flights %>% slice(n())`

`flights %>% slice(5:n())`

`slice(flights,-(1:4))`

- `flights %>% slice_sample(n=5)`
- `flights %>% slice_sample(n=5, replace = TRUE)`
- **set seed(123)**
- `train_data <- flights %>% slice_sample(prop=0.8)`
- `train_data`
- `test_data <- flights %>% slice_sample(prop=0.2)`
- `test_data`

Question/Queries?

# Database and R

- R Data Import/Export (rio):
  - <https://cran.r-project.org/doc/manuals/r-release/R-data.html>
- R Data Import/Export (rio): Relational database
  - <https://cran.r-project.org/doc/manuals/r-release/R-data.html#Relational-databases>
- Why use a database?
  - There are limitations on the types of data that R handles well.

# Why use a database?

- Since all data being manipulated by R are resident in memory, and several copies of the data can be created during execution of a function, **R is not well suited to extremely large data sets.**
- Data objects that are more than a (few) hundred megabytes in size can cause R to run out of memory, particularly on a 32-bit operating system.
- **R does not easily support concurrent access to data.** That is, if more than one user is accessing, and perhaps updating, the same data, the changes made by one user will not be visible to the others.

# Why use a database?

- **R does support persistence of data**, in that you can save a data object or an entire worksheet from one session and restore it at the subsequent session, but the format of the stored data is specific to R and not easily manipulated by other systems.
- Database management systems (DBMSs) and, in particular, relational DBMSs (RDBMSs) *are* designed to do all of these things well.
- The sort of statistical applications for which DBMS might be used are to extract a 10% sample of the data, to cross-tabulate data to produce a multi-dimensional contingency table, and to extract data group by group from a database for separate analysis.

# R interface package for Database:

- There are several packages available on CRAN to help R communicate with DBMSs. They provide different levels of abstraction.
- Some provide means to copy whole data frames to and from databases.
- All have functions to select data within the database via SQL queries, and to retrieve the result as a whole as a data frame or in pieces (usually as groups of rows).
- All **except RODBC** are tied to one DBMS, but there has been a proposal for a unified 'front-end' package **DBI** (<https://developer.r-project.org/db/>) in conjunction with a 'back-end', the most developed of which is **RMySQL**.

# Packages using DBI

- Package RMySQL on CRAN provides an interface to the MySQL database system (see <https://www.mysql.com> and Dubois, 2000) or its fork MariaDB (see <https://mariadb.org/>).
- The description here applies to versions 0.5-0 and later: earlier versions had a substantially different interface.
- The current version requires the DBI package, and this description will apply with minor changes to all the other back-ends to DBI.



# Packages using DBI: MySQL

- MySQL exists on Unix/Linux/macOS and Windows: there is a 'Community Edition' released under GPL but commercial licenses are also available.
- MySQL was originally a 'light and lean' database. (It preserves the case of names where the operating file system is case-sensitive, so not on Windows.)
- The call **dbDriver("MySQL")** returns a database connection manager object, and then a call to **dbConnect** opens a database connection which can subsequently be closed by a call to the generic function **dbDisconnect**.
- Use **dbDriver("Oracle")**, **dbDriver("PostgreSQL")** or **dbDriver("SQLite")** with those DBMSs and packages **ROracle**, **RPostgreSQL** or **RSQLite** respectively.

# Packages using DBI: MySQL

- SQL queries can be sent by either **dbSendQuery** or **dbGetQuery**.
- **dbGetquery** sends the query and retrieves the results as a data frame.
- **dbSendQuery** sends the query and returns an object of class inheriting from "**DBIResult**" which can be used to retrieve the results, and subsequently used in a call to **dbClearResult** to remove the result.
- Function **fetch** is used to retrieve some or all of the rows in the query result, as a list.
- The function **dbHasCompleted** indicates if all the rows have been fetched, and **dbGetRowCount** returns the number of rows in the result.

# Packages using DBI: MySQL

##install “RMySQL” package, if needed!

##load “RMySQL” package

- library(RMySQL) # will load DBI as well;

## open a connection to a MySQL database

- con <- dbConnect(dbDriver("MySQL"), dbname = "test")

## list the tables in the database

- dbListTables(con)

## load a data frame into the database, deleting any existing copy

- data(USArrests)
- dbWriteTable(con, "arrests", USArrests, overwrite = TRUE)
- dbListTables(con)

# Packages using DBI: MySQL

- `## get the whole table`
- `> dbReadTable(con, "arrests")`

|            | Murder | Assault | UrbanPop | Rape |
|------------|--------|---------|----------|------|
| • Alabama  | 13.2   | 236     | 58       | 21.2 |
| • Alaska   | 10.0   | 263     | 48       | 44.5 |
| • Arizona  | 8.1    | 294     | 80       | 31.0 |
| • Arkansas | 8.8    | 190     | 50       | 19.5 |
| • ...      |        |         |          |      |

# Packages using DBI: MySQL

## Select from the loaded table

```
dbGetQuery(con, paste("select row_names, Murder from arrests",
 "where Rape > 30 order by Murder"))
```

## Remove table

```
dbRemoveTable(con, "arrests")
```

## disconnect from database

```
dbDisconnect(con)
```

# Package RODB

- Package RODB on CRAN provides an interface to database sources supporting an Microsoft Open Database Connectivity (ODBC) interface.
- This is very widely available, and allows the same R code to access different database systems.
- RODB runs on Unix/Linux, Windows and macOS, and almost all database systems provide support for ODBC.
- It has been tested Microsoft SQL Server, Access, MySQL, PostgreSQL, Oracle and IBM DB2 on Windows and MySQL, MariaDB, Oracle, PostgreSQL and SQLite on Linux.

# Package RODB

- ODBC is a client-server system, and it has been successfully connected to a DBMS running on a Unix server from a Windows client, and vice versa.
- On Windows ODBC support is part of the OS.
- On Unix/Linux you will need an ODBC Driver Manager such as unixODBC (<http://www.unixODBC.org>) or iODBC (<http://www.iODBC.org>: this is pre-installed in macOS) and an installed driver for your database system.
- Windows provides drivers not just for DBMSs but also for Excel (.xls) spreadsheets, DBase (.dbf) files and even text files.

# Package RODBC

- Many simultaneous connections are possible.
- A connection is opened by a call to **odbcConnect** or **odbcDriverConnect**, which returns a handle used for subsequent access to the database.
- Printing a connection will provide some details of the ODBC connection, and calling **odbcGetInfo** will give details on the client and server.
- A connection is closed by a call to `close` or **odbcClose**, and also (with a warning) when not R object refers to it and at the end of an R session.



# Package RODB

- Details of the tables on a connection can be found using **sqlTables**.
- Function **sqlSave** copies an R data frame to a table in the database, and **sqlFetch** copies a table in the database to an R data frame.
- An SQL query can be sent to the database by a call to **sqlQuery**. This returns the result in an R data frame. (**sqlCopy** sends a query to the database and saves the result as a table in the database.)
- A finer level of control is attained by first calling **odbcQuery** and then **sqlGetResults** to fetch the results. The latter can be used within a loop to retrieve a limited number of rows at a time, as can function **sqlFetchMore**.

# Package RODB:

## Example using PostgreSQL

```
install RODB package, if needed
```

```
#Load RODB package
```

- library(RODB)

```
tell it to map names to l/case
```

- channel <- odbcConnect("testdb", uid="ripley", case="tolower")

```
load a data frame into the database
```

- data(USArrests)

- sqlSave(channel, USArrests, rownames = "state", addPK = TRUE)

- rm(USArrests)

```
list the tables in the database
```

- sqlTables(channel)

# Package RODBC

```
list it
```

```
sqlFetch(channel, "USArrests", rownames = "state")
```

```
an SQL query, originally on one line
```

```
sqlQuery(channel, "select state, murder from USArrests
where rape > 30 order by murder")
```

```
remove the table
```

```
sqlDrop(channel, "USArrests")
```

```
close the connection
```

```
odbcClose(channel)
```

# Getting Microsoft Access data in R: (Self-Learning)

- <https://www.r-bloggers.com/2013/01/getting-access-data-into-r/>
- <https://leowong.ca/blog/connect-to-microsoft-access-database-via-r/>
- <https://cran.r-project.org/web/packages/RODBC/RODBC.pdf>
- Same arguments holds true for Excel and any text files as they can also be imported with RODBC package in R!

# dplyr and dbplyr packages for Database:

<https://r4ds.hadley.nz/databases.html>

- Setting up a client-server or cloud DBMS is out of the scope of this course, so we'll instead use an in-process DBMS that lives entirely in an R package: **duckdb**.
- `install.packages("duckdb")`
- Connecting to **duckdb** is particularly simple because the defaults create a temporary database that is deleted when you quit R.
- `library(duckdb)`
- `library(dbplyr)`    `#Also required for "db" related functions`

dplyr and dbplyr packages for Database:  
<https://r4ds.hadley.nz/databases.html>

#Making a connection:

- `con <- DBI::dbConnect(duckdb::duckdb())`

#Con is a new database so we need to add some tables there

- `DBI::dbWriteTable(con, "mpg", ggplot2::mpg)`
- `DBI::dbWriteTable(con, "diamonds", ggplot2::diamonds)`

#Check the list of tables in con with:

- `DBI::dbListTables(con)`

# dplyr and dbplyr packages for Database:

<https://r4ds.hadley.nz/databases.html>

- #What will happen with this?
- `con %>% DBI::dbReadTable("diamonds") %>% as_tibble()`

#You/we can use the “SQL” to get the same tibble as follows:

```
sql <- "
 SELECT carat, cut, clarity, color, price
 FROM diamonds
 WHERE price > 15000
"
```

#Print the sql object as tibble as follows:

```
as_tibble(DBI::dbGetQuery(con, sql))
```

# dbplyr basics:

#To use dbplyr, you must first use **tbl()** to create an object that represents a database table:

- `diamonds_db <- tbl(con, "diamonds")`      #load ggplot2 for this data
- `diamonds_db`

#This object is lazy; when you use dplyr verbs on it, dplyr doesn't do any work: it just records the sequence of operations that you want to perform and only performs them when needed.

- `big_diamonds_db <- diamonds_db %>% filter(price > 1500) %>%  
 select(carat:clarity, price)`
- `big_diamonds_db`



# dbplyr basics:

- You can tell this object represents a database query because it prints the DBMS name at the top, and while it tells you the number of columns, it typically doesn't know the number of rows.
- This is because finding the total number of rows usually requires executing the complete query!
- `big_diamonds_db %>% show_query()` **#See SQL code generated by dbplyr**
- `<SQL>`
- `SELECT "carat", "cut", "color", "clarity", "price"`
- `FROM "diamonds"`
- `WHERE ("price" > 1500.0)`

# dbplyr basics:

- To get all the data back into R, you call `collect()`.
- Behind the scenes, this generates the SQL, calls `dbGetQuery()` to get the data, then turns the result into a tibble:
- `big_diamonds <- big_diamonds_db %>% collect()`
- `big_diamonds`

# dbplyr basics:

- Typically, you'll use dbplyr to select the data you want from the database, performing basic filtering and aggregation using the translations described earlier.
- Then, once you're ready to analyse the data with functions that are unique to R, you'll `collect()` the data to get an in-memory tibble, and continue your work with pure R code.

# What will happen?

- `dbplyr::copy_nycflights13(con)`
- `flights <- tbl(con, "flights")`
- `flights %>% show_query()`
- `planes <- tbl(con, "planes")`
- `planes %>% show_query()`

# What will happen?

- `flights %>% filter(dest == "IAH") %>% arrange(dep_delay) %>% show_query()`
- **WHERE and ORDER BY control which rows are included and how they are ordered**
- `flights %>% group_by(dest) %>% summarize(dep_delay = mean(dep_delay, na.rm = TRUE)) %>% show_query()`
- **GROUP BY converts the query to a summary, causing aggregation to happen**

# More here:

- Chapter 22: Databases
- R for Data Science, 2<sup>nd</sup> Edition
- <https://r4ds.hadley.nz/>

Question/Queries?

# Thank you!

@shitalbhandary